

CHAPTER V: PROCESS CONTROL IN UNIX

B. Boufama

UNIVERSITY OF WINDSOR

Unix Processes

An executing program \rightarrow a *process*.

In particular, every process in Unix has the followings :

- A unique process ID (PID)
- Some code : instructions that are being executed
- Some data : variables
- A stack : a form of memory where it is possible to push and pop.
- An environment : registers' contents, tables of open files,...

Unix starts as a single process, called *init*. The PID of *init* is 1.

The only way to create a new process in Unix, is to duplicate an existing one.

→ the process *init* is the ancestor of all subsequent processes.

In particular, process *init* never dies.

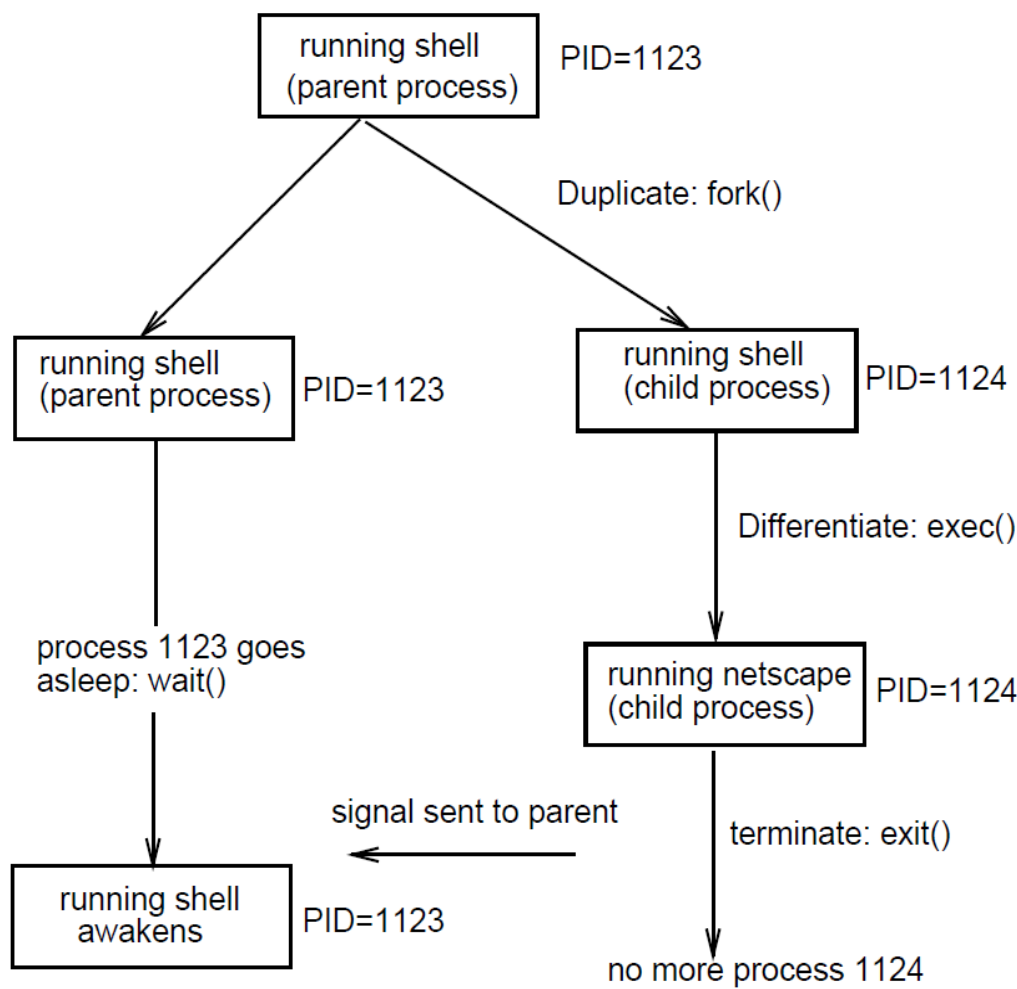
The creation or spawning of new processes is done with two system calls :

- *fork()*: duplicates the caller process
- *exec()*: replaces the caller process by a new one.

Example: Running a utility from a shell(*cs**h*).

The following steps are necessary

- *cs**h* *forks* a copy of itself.
- the child process *execs* the utility program, for example, *netscape*.
- the parent process waits for the termination(*exit*) of its child process by going asleep.
- when the child process terminates, a signal is sent to the parent process(the shell program *cs**h*). The latter wakes-up and becomes ready to accept the next command.



Creating a new Process: `fork()`

Synopsis: **`pid_t fork(void);`**

when succesful, the *fork()* system call :

- creates a copy of the caller (parent) process.
- returns the *PID* of the newly created process to the parent
- returns 0 to the new process (the child).

If not succesful, the *fork()* returns -1.

fork() is a strange system call : called by a single process but returns twice, to two different processes.

In particular, a child process has :

- its own unique *PID*,
- a different *PPID*,
- its own copy of the parent's data segment and file descriptors

fork() is primarily used in two situations :

1. A process wants to execute another program (Shells).
2. A process has a main task and when necessary creat a child to handle an operation(Servers).

Here is a simple example :

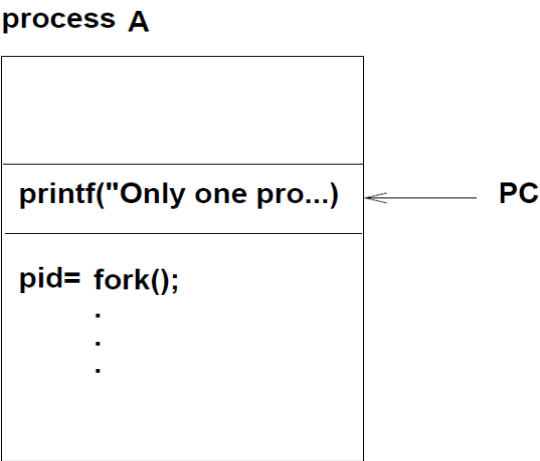
```
#include <unistd.h>

int main(int argc, char *argv[]){
    int pid;

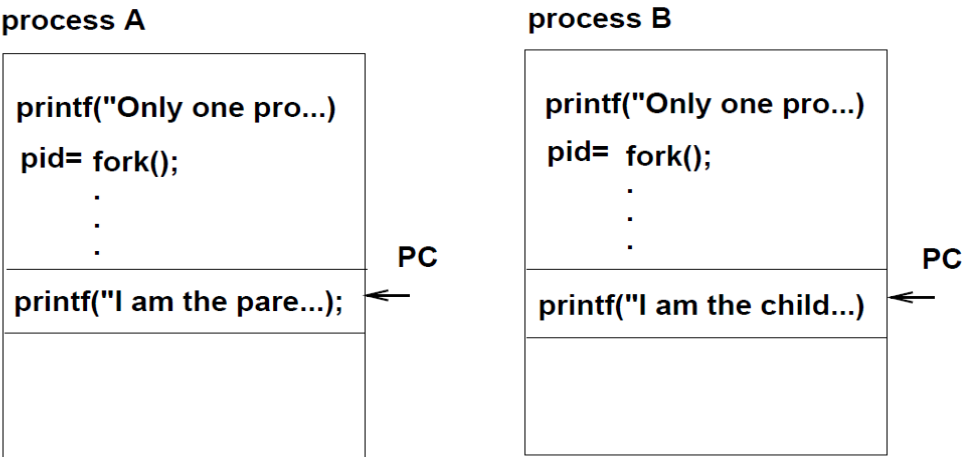
    printf("Only one process\n");
    pid = fork();

    if(pid == -1){
        perror("impossible to fork");
        exit(1);
    }
    if(pid > 0)
        printf("I am the parent, pid=%d\n", getpid());
    else
        if(pid == 0)
            printf("I am the child, pid=%d\n", getpid());

    exit(0);
}
```

After fork



Terminating a process: `exit()`

Synopsis: **`void exit(int status);`**

This call terminates a process and never returns

The *status* value is available to the parent process through the *wait()* system call.

When invoked by a process, the *exit()* system call:

- closes all the process's file descriptors
- frees the memory used by its code, data and stack
- sends a **SIGCHLD** signal to its parent and waits for the parent to accept its return code.

Waiting for a process: `wait()`

Synopsis: **`pid_t wait(int *status);`**

This call allows a parent process to wait for one of its children to terminate and to accept its child's termination code.

When called, *wait()* can

- block (suspend) the caller process, if all of its children are still running, or
- return immediately with a termination status of a child, if a child has terminated and is waiting for its termination to be accepted, or
- return immediately with an `error(-1)` if it does not have any child process.

Waiting for a specific process: `waitpid()`

Synopsis:

pid of a specific child not all

```
pid_t waitpid(pid_t pid, int *status, int options);
```

This call allows a parent process to wait for a specific child to terminate and to accept its child's termination code.

One interesting option is **WNOHANG** that causes the **waitpid()** to return immediately even if no child has exited.

Note: **wait(&status)** is equivalent to
waitpid(-1, &status, 0)
any child

When succesful, *wait()* returns the *pid* of the terminating child process.

The value in *status* is encoded as follow :

- if the rightmost byte of *status* is zero, then the leftmost byte contains the status returned by the child : a value between 0 and 255.




This represents a normal termination of the child process.

- if the rightmost byte of *status* is nonzero, then the rightmost 7 bits are equal to the *signal number*, that caused the process to terminate. The remaining bit of the rightmost byte is set to 1 if a core dump was produced by the child process.

Some bit-manipulation macros have been defined to deal with the value in the variable *status*(you need to include `< sys/wait.h >`).

- *WIFEXITED(status)*: true for normal child termination.
- *WEXITSTATUS(status)*: used only when *WIFEXITED(status)* is true, it returns the exit status as an integer(0-255).
- *WIFSIGNALED(status)*: true for abnormal child termination
- *WTERMSIG(status)*: used only when *WIFSIGNALED(status)* is true, it returns the signal number that caused the abnormal child death.

status passed to wait is 4 bytes which is 32 bits



8 bits

Orphan and zombie Processes

A process that terminates does not leave the system before its parent accepts its return.

There are 2 interesting situations, when :

1. a parent exits(for example, the parent has been killed prematurely) while its children are still alive.
→ the children become *orphans*.

Because somebody must accept their return codes, the kernel simply changes their *PPID* to 1.

→ *orphan* processes are systematically adopted by the process *init* (PID of init is 1).

In particular, *init* accepts all its children returns.

2. a live parent never makes the system call *wait()*. the children become *zombies* and remain in the system's process table waiting for the acceptance of their return. However, they loose their ressources (data, code, stack...).

Because the system's process table has a fixed-size, too many zombie processes can require the intervebtion of the system administrator.

Example :

Below is a C program, called *zombie.c*, to create a zombie.


```
int main(int argc, char *argv[]){
    int pid;

    pid = fork();
    if (pid){ // means pid !=0
        printf("parent process, pid=%d\n", getpid());
        while(1)
            sleep(5);
    }
    printf("child process, pid=%d\n", getpid());
    exit(0);
}
```

1- run in background the program:> ./zombie &

On the screen you will get:

I am the child, pid=12357

I am the parent, pid=12356

2- look for the processes:> ps -ef | grep 12356

On the screen you will get:

boufama 12357 12356 0 0:00 < *defunct* >

boufama 12356 20142 0 20:01:44 pts/7 0:00 ./zombie

Differentiating a process: `exec()`

The *exec()* family of system calls allows a process to replace its current code, data and stack with those of another program.

Synopsis :

- `int execl(const char *path, [const char *argi,]+ NULL)`
- `int execlp(const char *path, [const char *argi,]+ NULL)`
- `int execv(const char *path, const char *argv[])`
- `int execvp(const char *path, const char *argv[])`

where $i = 0, \dots, n$ and ⁺ means one or more times.

The difference between these 4 system calls has to do with syntax.

execl() and *execv()* require the whole pathname of the executable program to be supplied.

execlp() and *execvp()* use the variable **\$PATH** to find the program.

In particular :

- A successful call to *exec()* never returns.
- *exec()* returns -1 if not successful.
- For both *execl()* and *execlp()* *arg₀* must be the name of the program.
- For both *execv()* and *execvp()* *arg[0]* must be the name of the program.

Changing directories: `chdir()`

A child process inherits its current working directory from its parent.

Each process can change its working directory using *chdir()*.

Synopsis :

```
int chdir(const char * pathName);
```

chdir() returns 0 if successful -1 otherwise.

It fails if the specified path name does not exist or if the process does not have execute permission from the directory.