

CHAPTER VII: PIPES

B. Boufama

UNIVERSITY OF WINDSOR

Unnamed Pipes

Unnamed Pipes, known as *pipe*, are a mechanism for interprocess communication.

In particular, they are used by shells to connect one utility's standard output with the standard input of another utility.

Example: `ps -ef | grep Viber | wc -l`

Pipes are the oldest form of Unix IPC.

Pipes have two limitations, they :

- are half-duplex: data flows in one direction only
- can be used only between processes that have a common ancestor. Typically, a process creates a pipe, forks then, uses the pipe to exchange information with its child.


The *pipe()* System call

Synopsis : `int pipe(int fd[2])`

returns 0 when successful and -1 otherwise

Doesnt return but sets read to fd[0] and write to fd[1]

pipe() creates a pipe and ~~returns two~~ file descriptors, fd[0] and fd[1], where fd[0] is open for reading and fd[1] is open for writing.



→ a pipe is a one-way communication channel between two related processes.

When reading from or writing to a pipe, the following rules apply :

- If a process reads from a pipe whose write end has been closed, after all data has been read, `read()` returns 0 (end-of-file).
- If a process reads from an empty pipe whose write end is still open, it sleeps until some input becomes available.
- If a process tries to read from a pipe more bytes than are present, `read()` reads all available bytes and returns the number of bytes read.
- If a process writes to a pipe whose read end has been closed, the write operation fails and the writer process receives a SIGPIPE.

Note : In case of multiple processes writing to the same pipe, a *write* of up to `PIPE_BUF` bytes is guaranteed to be atomic

→ data from different writer processes will not be interleaved.

Not interleaved - Other process will wait for one process to fully write(it happens in a single step as write is atomic).

PIPE_BUF - max these many characters only can be written atomically else it will be more than 1 step. eg if PIPE_BUF is 1000 then only 1000 characters can be written atomically

Example 2: child is a writer, parent is a reader

```
#include <stdio.h>
#include <unistd.h>

void child(int *);
void parent(int *);

int main(int argc, char *argv[]){
    int fd[2];

    if(pipe(fd) == -1)
        exit(1);

    if(fork() == 0)
        child(fd);
    else
        parent(fd);
    exit(0);
}
```

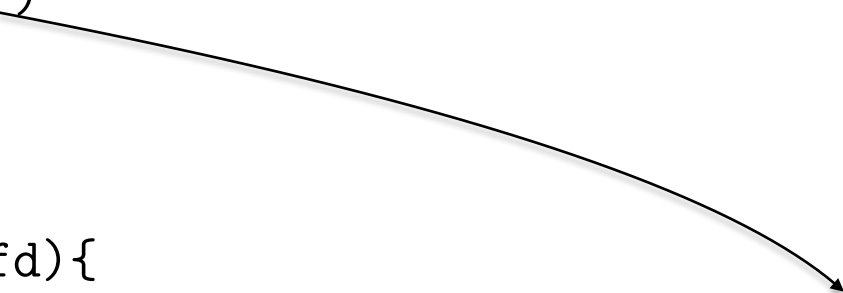
```
void parent(int *fd){
    char ch;

    close(fd[1]);
    printf("Child has sent the message:\n");
    do{
        read(fd[0], &ch, 1);
        printf("%c", ch);
        if(ch == '\n')
            break;
    }while(1);
}

void child(int *fd){
    char message[255]="Hello, here is my data...\n";

    close(fd[0]);

    write(fd[1], message, 26);
}
```



Note that two pipes are necessary in order to get a bidirectional communication.

The system call: `dup2()`

Synopsis : `int dup2(int fd, int fd2);`

The `dup2()` function causes the file descriptor `fd2` to refer to the same file as `fd`. The `fd` argument is a file descriptor referring to an open file, and `fd2` is a non-negative integer less than the current value for the maximum number of open file descriptors allowed the calling process.

If `fd2` already refers to an open file, not `fd`, it is closed first. If `fd2` refers to `fd`, or if `fd` is not a valid open file descriptor, `fd2` will not be closed first.

If successful, `dup2()` returns a non-negative integer representing the file descriptor `fd2`. Otherwise, -1 is returned

0 means `stdin`(read from the console) but `dup2(fd,0)` will read the contents from file instead.

Example: `dup2(fd, 0);`

→ reading from `stdin` will mean reading from the file whose descriptor is `fd`.

Example 3: implementing the shell pipe mechanism

```
int main(int argc, char *argv[]){
    int fd[2];

    if(pipe(fd)==-1) exit(1);
    if(fork() > 0)
        parent(fd, argv);
    else child(fd, argv);
}

void parent(int *fd, char *argv[]){ // A writer
    close(fd[0]); // closing the read end of pipe
    dup2(fd[1], 1); // 1 is the standard output
    // instead of writing stdout's contents to console, dup2 will write it to pipes write end
    close(fd[1]); // close original file descriptor
    execlp(argv[1], argv[1], NULL);
}
// exec will output the result to console thru stdout, but after dup2 it will be written pipes write
// end thru stdout.

void child(int *fd, char *argv[]){ // A reader
    close(fd[1]); // closing the write end of pipe
    dup2(fd[0], 0); // 0 is the standard input
    // instead of reading stdin's contents from console, dup2 will make stdin to read it from pipes read end
    close(fd[0]); // close original file descriptor
    execlp(argv[2], argv[2], NULL);
}
```


Example 4: implementing the shell pipe mechanism

```
int main(int argc, char *argv[]){
    int fd1[2], fd2[2], fd3[2], fd4[2];
    char turn='T';

    printf("This is a 2-player game with a referee\n");
    pipe(fd1);
    pipe(fd2);
    if(!fork())
        player("TOT0", fd1, fd2);

    close(fd1[0]);    // parent only write to pipe 1
    close(fd2[1]);    // parent only reads from pipe 2

    pipe(fd3);
    pipe(fd4);
    if(!fork())
        player("TITI", fd3, fd4);

    close(fd3[0]);    // parent only write to pipe 3
    close(fd4[1]);    // parent only reads from pipe 4
```

```
while(1){  
    printf("\nReferee: TOT0 plays\n\n");  
    write(fd1[1], &turn, 1);  
    read(fd2[0], &turn, 1);  
  
    printf("\nReferee: TITI plays\n\n");  
    write(fd3[1], &turn, 1);  
    read(fd4[0], &turn, 1);  
}  
}
```

```
void player(char *s, int *fd1, int *fd2){
    int points=0;
    int dice;
    long int ss=0;
    char turn;

    while(1){
        read(fd1[0], &turn, 1);
        printf("%s: playing my dice\n", s);
        dice =(int) time(&ss)%10 + 1;
        printf("%s: got %d points\n", s, dice);
        points+=dice;
        printf("%s: Total so far %d\n\n", s, points);
        if(points >= 50){
            printf("%s: game over I won\n", s);
            kill(0, SIGTERM);
        }
        sleep(5);    // to slow down the execution
        write(fd2[1], &turn, 1);
    }
}
```

FIFOs or Named Pipes

FIFOs(First In First Out), sometimes called named pipes, offer the following advantages over pipes :

- They have a name that exists in the file system.
- They can be used by unrelated processes.
- They exist until explicitly deleted.

The system call *mkfifo()*:

int mkfifo(const char *path, mode_t mode)

mkfifo() return 0 if OK, -1 otherwise.

Creating a FIFO is similar to creating a file.

Example: `mkfifo("/tmp/channel.fif", "0755");`

Note: *mkfifo()* is actually a library function that invokes system call *mknod()*, to create special files.

Once a FIFO has been created, it can be treated as a file. In particular, the system calls *open()*, *close()*, *read()*, *write()* and *unlink()* (to delete a file) can be used on a FIFO.

By default, we have :

- Calling *open()* for read only blocks the caller until some other process opens the FIFO for writing.
- Calling *open()* for write only blocks the caller until some other process opens the FIFO for reading.
- If a process writes to a FIFO that no process has open for reading, the signal *SIGPIPE* will be generated.
- When the last writer for a FIFO closes the FIFO, an end-of-file will be generated for the reader.
- Like pipes, FIFOs are one-way communication channels.

Note : In case of multiple processes writing to the same pipe, a *write* of up to *PIPE_BUF* bytes is guaranteed to be atomic

Example 5: a client/server application where a server accepts data from clients using the FIFO /tmp/server.

```
#include <fcntl.h>
#include <stdio.h>          // This is the server

int main(int argc, char *argv[]){
    int fd;
    char ch;

    unlink("/tmp/server"); // delete it if it exists
    if(mkfifo("/tmp/server", 0777)!=0)
        exit(1);
    chmod("/tmp/server", 0777);
    while(1){
        fprintf(stderr, "Waiting for a client\n");
        fd = open("/tmp/server", O_RDONLY);
        fprintf(stderr, "Got a client: ");
        while(read(fd, &ch, 1) == 1)
            fprintf(stderr, "%c", ch);
    }
}
```

```
#include <fcntl.h>                // This is the client
#include <stdio.h>

int main(int argc, char *argv[]){
    int fd;
    char ch;

    while((fd=open("/tmp/server", O_WRONLY))==-1){
        fprintf(stderr, "trying to connect\n");
        sleep(1);
    }

    printf("Connected: type in data to be sent\n");

    while((ch=getchar()) != -1) // -1 is CTR-D
        write(fd, &ch, 1);

    close(fd);
}
```

Example 6: a client/server application where a server creates a child for each client. Then, the child creates a separate FIFO to send data to the client.

```
#include <wait.h>
#include <fcntl.h>           // This is the server
#include <stdio.h>

void child(pid_t client);

int main(int argc, char *argv[]){
    int fd, status;
    char ch;
    pid_t pid;

    unlink("/tmp/server");
    if(mkfifo("/tmp/server", 0777)){
        perror("main");
        exit(1);
    }
    chmod("/tmp/server", 0777);
```



```
while(1){
    fprintf(stderr, "Waiting for a client\n");
    fd = open("/tmp/server", O_RDONLY);
    fprintf(stderr, "Got a client: ");
    read(fd, &pid, sizeof(pid_t));
    fprintf(stderr, "%ld\n", pid);
    if(fork()==0){
        close(fd);
        child(pid);
    }else
        waitpid(0, &status, WNOHANG);
}
```

```
void child(pid_t pid){
    char fifoName[100];
    char newline='\n';
    int fd, i;

    sprintf(fifoName, "/tmp/fifo%d", pid);
    mkfifo(fifoName, 0777);
    chmod(fifoName, 0777);
    fd = open(fifoName, O_WRONLY)

    for(i=0; i < 10; i++){
        write(fd, fifoName, strlen(fifoName));
        write(fd, &newline, 1);
    }
    close(fd);
    unlink(fifoName);
    exit(0);
}
```

```
int main(int argc, char *argv[]) { // Client
    char ch, fifoName[50];
    int fd;
    pid_t pid;

    while((fd=open("/tmp/server", O_WRONLY))==-1){
        fprintf(stderr, "trying to connect\n");
        sleep(1);
    }
    pid = getpid();
    write(fd, &pid, sizeof(pid_t));
    close(fd);

    sprintf(fifoName, "/tmp/fifo%d", pid);
    sleep(1);

    fd = open(fifoName, O_RDONLY);
    while(read(fd, &ch, 1) == 1)
        fprintf(stderr, "%c", ch);
    close(fd);
}
```